

# Analysis 1.1

January 13, 2019

## 1 Data Science Popularization

Dataset: DSP Master version-1

Relevant Information: The dataset contains instances of reporters from various fields of expertise with relevant information of the reporters which are sparse, repetitive and often missing.

Number of instances : 3316

Number of attributes : 16

Attributes : Country, Province, Category, Newspaper, Newspaper Website, Magazine, Magazine Website, Reporter Name, Reporter Email Reporter Phone, Reporter Twitter, Reporter Linkedin, Technology Health&Medical, Environment, Others

Missing Attribute values : Abundant

### 1.1 The journey with the master dataset :

1. Primarily we preprocessed the master dataset by dropping the duplicates
2. To classify names with respective gender scraped the possible Gender attributed names from various sites (viz, <https://www.familyeducation.com/baby-names/browse-names/first-name/>) and stored it in a github repo for reference (<https://github.com/adjunctexcist01/namesdb>).
3. To verify the data, we checked the user twitter profiles wherever applicable
4. We filled the missing Location values from the data collected from twitter. The resulting dataset had 2972 datapoints
5. Next, we scraped google scholar for the reporter names to extract their Interests, Affiliation, Citation Count, Titles of their publications, and 10 publications by each reporter. The resulting dataset had 600 datapoints
6. We preprocessed the textual data by removing html tags and stop-words.

**Problems we faced:** - Many of the scraped names were of unisex in nature and it was hard to classify it under a particular gender - We saw a huge drop in valid datapoints as ~2372 datapoints threw exception while scraping from scholarly

## 2 Enter the revised dataset

Dataset: Final-DSP

Relevant Information: The dataset is a final version of the dataset provided, which is truncated and filled with information specific to analysis.

Missing values: Minimal

## 2.1 Importing relevant libraries

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")
warnings.filterwarnings("ignore", category=DeprecationWarning)

import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer, TfidfVectorizer, CountVec

from sklearn import metrics
from sklearn.manifold import TSNE
from nltk.stem.porter import PorterStemmer
from sklearn.preprocessing import Normalizer

import re
from nltk.corpus import stopwords
from nltk.stem.wordnet import WordNetLemmatizer
from nltk.corpus import wordnet

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
from tqdm import tqdm
import os
import bokeh.plotting as bp
from bokeh.models import HoverTool, BoxSelectTool
from bokeh.plotting import figure, show, output_notebook, reset_output
from bokeh.palettes import d3
import bokeh.models as bmo
from bokeh.io import save, output_file
from sklearn.decomposition import TruncatedSVD
from sklearn.cluster import MiniBatchKMeans, KMeans
from sklearn.metrics import silhouette_score

from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import Normalizer

In [3]: df = df = pd.read_excel("/home/ae/Documents/Work/Publications/Final-DSP.xlsx") #data f
```

## 2.2 Shape of the dataset

```
In [4]: df.shape
```

```
Out[4]: (600, 31)
```

```
In [5]: df.columns
```

```
Out[5]: Index(['Category', 'Reporter Name', 'Technology', 'Health&Medical',  
              'Environment', 'Others', 'Country', 'Province', 'Newspaper',  
              'NewspaperWebsite', 'Magazine ', 'Magazine Website', 'ReporterEmail',  
              'Reporter Phone', 'ReporterTwitter', 'ReporterLinkedin', 'Gender',  
              'Twitter Link', 'Location', 'User Url', 'Bio', 'Affiliation', 'Citedby',  
              'Email', 'Interests', 'Publication Titles', 'CleanedText', 'Cleaned',  
              'Name', 'Publications', 'Cleaned Publications'],  
             dtype='object')
```

**NB :** - CleanedText and Cleaned Publications are Publication Titles and Publications respectively, which have undergone stop word, punctuation and html tag removal - Reporter Name is name from the provided Dataset and Name is the profile names resulted from the twitter scrape - Country is the location provided in the initial file and Location is derived from the twitter scrape - Bio is the twitter bio of individuals resulted from the scrape the importance of which seems negligible

### 2.3 Comparing the number of missing values between the provided and scraped data (location)

```
In [6]: df["Country"].describe()
```

```
Out[6]: count      583  
       unique      42  
       top        USA  
       freq       270  
       Name: Country, dtype: object
```

```
In [7]: #data entries having no known value for countries  
       df["Country"].isna().sum()
```

```
Out[7]: 17
```

```
In [8]: df["Location"].describe()
```

```
Out[8]: count      585  
       unique     141  
       top        USA  
       freq       237  
       Name: Location, dtype: object
```

```
In [9]: #data entries having no known value for their present location  
       df["Location"].isna().sum()
```

```
Out[9]: 15
```

## Inference

- Since there is a marginal difference between the provided location and the scraped ones, we are choosing the provided ones(Country) due to it's refined formatting, for analysis.


## 2.4 EDA on the final Dataset

### 2.4.1 Ratio of male/female reporters

```
In [10]: #count of male(M) and female(F) reporters  
df["Gender"].value_counts()
```

```
Out[10]: M    373  
        F    169  
        Name: Gender, dtype: int64
```

```
In [11]: sns.countplot(x="Gender", data=df).set_title("Total Gender Countplot")  
plt.show()
```



### Observation:


- For the final dataset, we see the count of male reporters is more than twice the number of female reporters. (For the case of successful gender label)

### 2.4.2 Reporters on the various domains

The data includes categories of reporting as : 1. Technology 2. Environment 3. Health 4. Others  
In the analysis below, we can compare the popularity of each domain among reporters.

```
In [12]: x = ['Technology', 'Health', 'Environment', 'Others']
y = [df['Technology'].sum(), df['Health&Medical'].sum(), df['Environment'].sum(), df['Others']]
barlist = plt.bar(x, y)
barlist[0].set_color("r")

plt.show()
```




From Observing the graph it is clear that the most of them are reporting on random or uncategorized issues.

### 2.4.3 Top 10 countries with highest number of reporters

```
In [13]: df.groupby(['Country'])["Reporter Name"].count().sort_values(ascending = False).head()


Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f844f153630>
```



USA has the highest number of reporters with around ~45% of the share in the total data.

#### 2.4.4 Gender ratio in the top 5 countries with the maximux number of reporters

```
In [14]: w=0.35
         ind = np.arange(5)
         male = df[(df["Gender"]=="M")].groupby(["Country"])["Reporter Name"].count().sort_val
         female = df[(df["Gender"]=="F")].groupby(["Country"])["Reporter Name"].count().sort_val
         plt.bar(ind, male, w, label="Male")
         plt.bar(ind+w, female, w, label="Female")
         plt.xticks(ind+w/2, male.keys())
         plt.legend()
         plt.grid()
         plt.show()
```




Most of the countries have female:male ratio as less than half.

#### 2.4.5 Top 5 countries with the highest number of reports in the Technology domain

```
In [15]: df[df["Technology"]==1].groupby(['Country'])["Reporter Name"].count().sort_values(ascending=True)
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x7f844f0ca128>
```



Top Countries in technology reporting are United Kingdom, Canada, Australia , USA(surprisingly, it ranks 4th despite having maximum reporters in the world) .


(The actual count of reporter for each country is significantly less as compared to the entire data, thus this is not a strong conclusion)

#### 2.4.6 Top 5 countries with the highest number of reports in the Health and Medicine domain

```
In [16]: df[df["Health&Medical"]==1].groupby(['Country'])["Reporter Name"].count().sort_values
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7f844efc5438>
```






#### 2.4.7 Top 5 countries with the highest number of reports in the Environment domain

```
In [17]: df[df["Environment"]==1].groupby(['Country'])["Reporter Name"].count().sort_values(ascending=True)
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f844ef4a470>
```



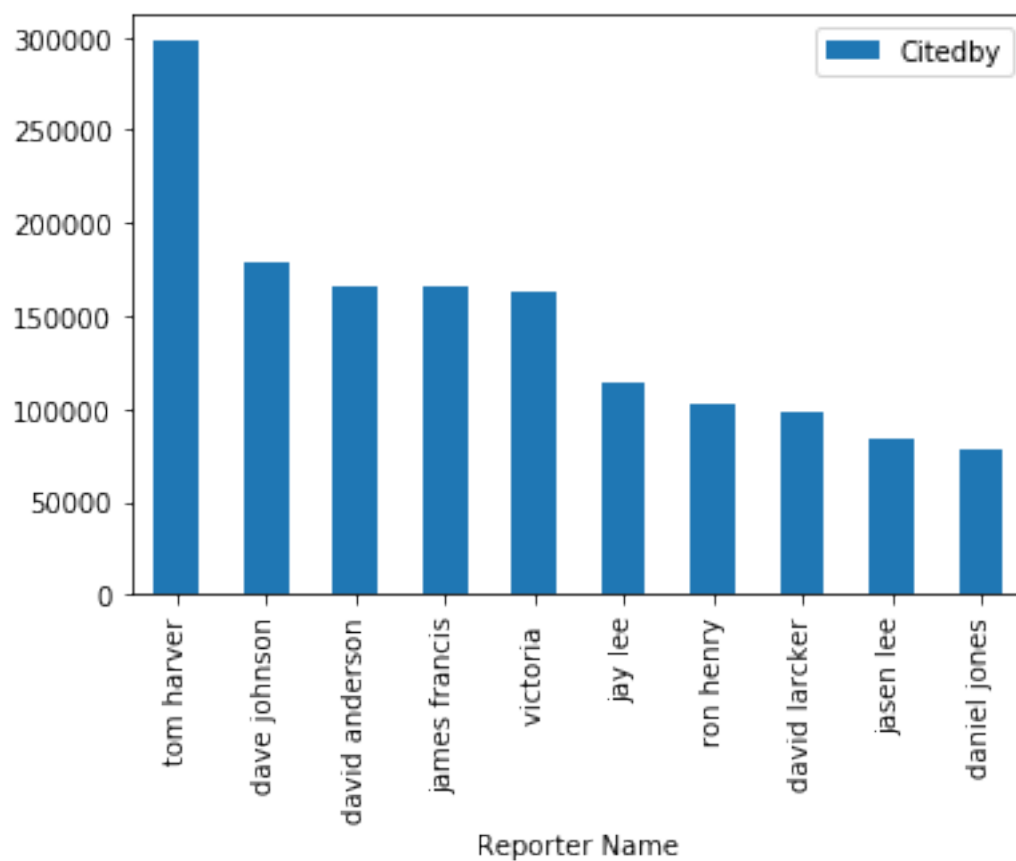
The count of reporters in Environment domain is low, thus we can't really confirm a pattern.

#### 2.4.8 People getting cited the most and their respective countries

**Most Active Reporters** On the basis of number of citations on google scholar authored by a reporter , we estimated the activity of a reporter and how frequently they indulge .


```
In [18]: df[(df["Citedby"].notnull())][["Reporter Name", "Citedby", "Country"]].sort_values(by
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f844ef317b8>
```



```
In [19]: df[(df["Citedby"].notnull())][["Reporter Name", "Citedby", "Country"]].sort_values(by
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x7f844ee3b3c8>
```



Tom Harver being the reporter with highest citations (close to 0.3 million)

Further, we have grouped them on the basis of country, to deduce the country most active in reporting. Top 3 actively reporting countries being:

1. United States of America
2. Canada
3. South Africa

## 2.5 Analysing the Textual data

### 2.5.1 Final preprocessing of Publication Text for exceptions

```
In [20]: #function to remove stopwords
def secondClean(text):
    cleaned = firstClean(text)
    filtered_words = []
    for x in cleaned.split():
        if x not in stop_words:
            filtered_words.append(x)
    return " ".join(filtered_words).strip()

#conversion to processable string
```

```

def toString(x):
    try:
        return str(x)
    except:
        return 0

#cleaning html tags from text
def firstClean(text):
    return re.sub(r"[^A-Za-z\u00C0-\u00D6\u00D8-\u00f6\u00f8-\u00ff\s]", "", text)
stop_words = stopwords.words('english')

In [21]: df['Cleaned Publications'] = df['Cleaned Publications'].apply(lambda x:toString(x))
In [22]: df['Cleaned Publications'] = df['Cleaned Publications'].apply(lambda x:secondClean(x))
In [23]: #tfidf used to find relevant words in text
vectorizer = TfidfVectorizer(min_df=5, analyzer='word', ngram_range=(1, 2), stop_words=stop_words)
vz = vectorizer.fit_transform(list(df['Cleaned Publications']))

vz.shape

Out[23]: (600, 8966)

In [24]: tfidf = dict(zip(vectorizer.get_feature_names(),vectorizer.idf_))
In [25]: tfidf = pd.DataFrame(columns=['tfidf']).from_dict(dict(tfidf), orient='index')
In [26]: tfidf.columns = ['tfidf']

```

## 2.5.2 Visualizing the textual data using Word Cloud

```

In [27]: #higher tfidf score, more relevant words is to the context
tfidf.head()

Out[27]:
          tfidf
abandoned    5.452685
abdominal    5.096010
aberrant      5.452685
aberrations  5.452685
abilities     4.508223

In [28]: from wordcloud import WordCloud

def plot_word_cloud(terms):
    text = terms.index
    text = ' '.join(list(text))
    wordcloud = WordCloud(max_font_size=40).generate(text)
    plt.figure(figsize=(25, 25))
    plt.imshow(wordcloud, interpolation="bilinear")
    plt.axis("off")
    plt.show()

```

```
In [29]: plot_word_cloud(tfidf.sort_values(by=['tfidf'], ascending=False).head(40))
```



**Inference:** This shows the TF-IDF cloud words based on the text of publications for each records, it is interesting to see that even though the data is dominated by tech reporters, TF-IDF shows words which are related to Health domains - cell - dna - renal - microbiota are some of such words

### 2.5.3 Clustering Publication Text

Purpose is to find similar reporters based on their publications through Clustering.

```
In [30]: # Single value reduction
from sklearn.decomposition import TruncatedSVD
svd = TruncatedSVD(n_components=50, random_state=0)
svd_tfidf = svd.fit_transform(vz)

svd_tfidf.shape
```

```
Out[30]: (600, 50)
```

```
In [31]: # T-distributed Stochastic Neighbor Embedding used. This helps in high level visualization
from sklearn.manifold import TSNE
tsne_model = TSNE(n_components=2, perplexity=10, verbose=1, random_state=0, n_iter=2000)
tsne_tfidf = tsne_model.fit_transform(svd_tfidf)
print(tsne_tfidf.shape)
tsne_tfidf_df = pd.DataFrame(tsne_tfidf)
tsne_tfidf_df.columns = ['x', 'y']
tsne_tfidf_df['Category'] = df['Category']
tsne_tfidf_df['ReporterName'] = df['Reporter Name']
```

```
[t-SNE] Computing 31 nearest neighbors...
[t-SNE] Indexed 600 samples in 0.004s...
[t-SNE] Computed neighbors for 600 samples in 0.045s...
[t-SNE] Computed conditional probabilities for sample 600 / 600
[t-SNE] Mean sigma: 0.158728
[t-SNE] KL divergence after 250 iterations with early exaggeration: 79.166100
[t-SNE] KL divergence after 2000 iterations: 0.953784
(600, 2)
```

```
In [33]: output_notebook()
plot_tfidf = bp.figure(plot_width=700, plot_height=600, title="tf-idf clustering of t
    tools="pan,wheel_zoom,box_zoom,reset,hover,previewsave",
    x_axis_type=None, y_axis_type=None, min_border=1)

palette = d3['Category10'][len(tsne_tfidf_df['Category'].unique())]
color_map = bmo.CategoricalColorMapper(factors=tsne_tfidf_df['Category'].map(str).uni

plot_tfidf.scatter(x='x', y='y', color={'field': 'Category', 'transform': color_map},
    legend='Category', source=tsne_tfidf_df)
hover = plot_tfidf.select(dict(type=HoverTool))
hover.tooltips={"Reporter Name": "@ReporterName", "Category": "@Category"}

show(plot_tfidf)
```

Inference: - Nearest Neighbour Clustering based on TSNE shows that there are small clusters scattered all over the place, hence reporters are fairly diverse in their style of publications. - The clusters are really small, thus it doesn't make sense to analyze each and every small cluster

## Distorsion and Silhouette Score on Publication

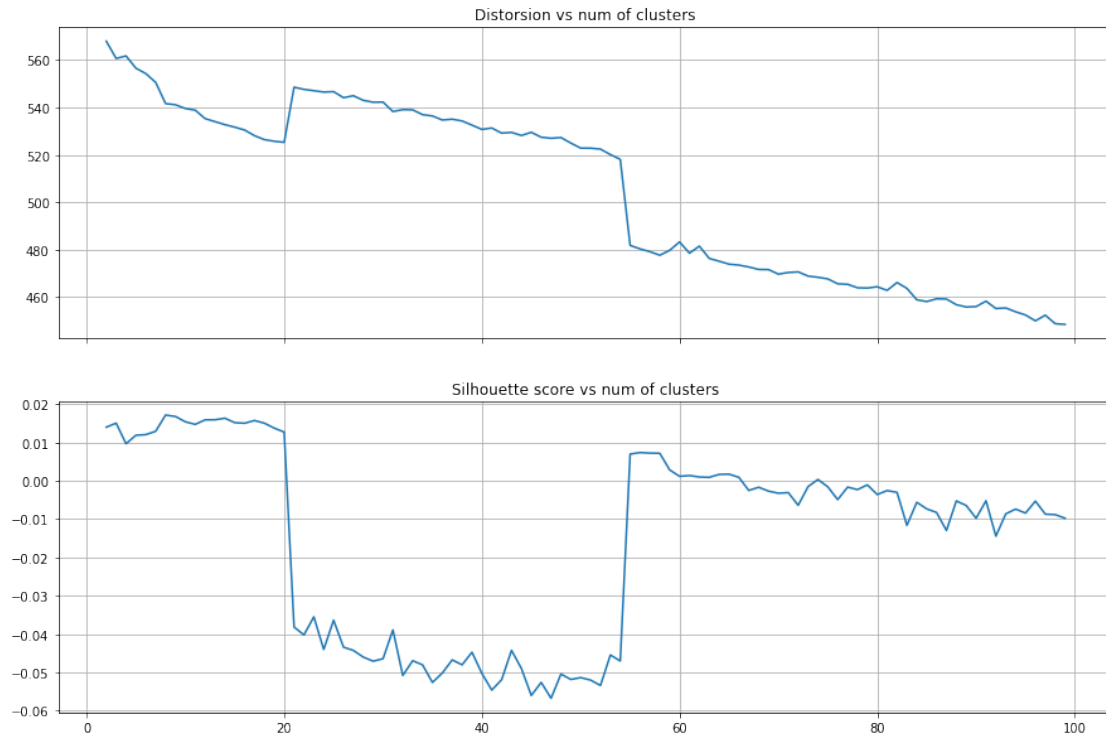
```
In [34]: from tqdm import tqdm_notebook
distorsions = []
sil_scores = []
k_max = 100
for k in tqdm_notebook(range(2, k_max)):
    kmeans_model = MiniBatchKMeans(n_clusters=k, init='k-means++', n_init=1, random_s
        init_size=1000, verbose=False, max_iter=1000)
    kmeans_model.fit(vz)
    sil_score = silhouette_score(vz, kmeans_model.labels_)
    sil_scores.append(sil_score)
    distorsions.append(kmeans_model.inertia_)

HBox(children=(IntProgress(value=0, max=98), HTML(value='')))
```

```
In [35]: f, (ax1, ax2) = plt.subplots(2, 1, sharex=True, figsize=(15, 10))
```

```
ax1.plot(range(2, k_max), distortions)
ax1.set_title('Distorsion vs num of clusters')
ax1.grid(True)

ax2.plot(range(2, k_max), sil_scores)
ax2.set_title('Silhouette score vs num of clusters')
ax2.grid(True)
```



Above graph shows Error to accuracy data points for each K in K-means till 100, there is an abnormality after 20 till around ~50 we dont know the reason of such abnormality, we also dont have an elbow curve to determine which K would be more suited Since the error is still decreasing even after we reach k=100. We cannot chose a large K because we dont have enough data for it cluster properly.if we choose k=100 it will select 6 datapoints in for each clusters (total records ~600). and analyzing those 100 clusters would not be feasible

## 2.6 Clustering on Affiliation

The Next interesting attribute to cluster on is on affiliation to find similarities, It includes Names of Universities, State etc

### Preprocessing Affiliations

```
In [36]: df['Cleaned Affiliations'] = df['Affiliation'].apply(lambda x:secondClean(x))
```



```

In [37]: vectorizerAff = TfidfVectorizer(min_df=5, analyzer='word', ngram_range=(1, 2), stop_w
vz1 = vectorizerAff.fit_transform(list(df['Cleaned Affiliations']))
df.reset_index(inplace =True)
vz1.shape

Out[37]: (600, 98)

In [38]: #single value reduction
svdAff = TruncatedSVD(n_components=50, random_state=0)
svd_tfidfAff = svdAff.fit_transform(vz1)
svd_tfidfAff.shape

Out[38]: (600, 50)

In [39]: tsne_modelAff = TSNE(n_components=2,perplexity=10, verbose=1, random_state=0, n_iter=
tsne_tfidfAff = tsne_modelAff.fit_transform(svd_tfidfAff)
print(tsne_tfidfAff.shape)
tsne_tfidf_dfAff = pd.DataFrame(tsne_tfidfAff)
tsne_tfidf_dfAff.columns = ['x', 'y']
tsne_tfidf_dfAff['Category'] = df['Category']
tsne_tfidf_dfAff['ReporterName'] = df['Reporter Name']
tsne_tfidf_dfAff['Affiliations'] = df['Cleaned Affiliations']

[t-SNE] Computing 31 nearest neighbors...
[t-SNE] Indexed 600 samples in 0.002s...
[t-SNE] Computed neighbors for 600 samples in 0.036s...
[t-SNE] Computed conditional probabilities for sample 600 / 600
[t-SNE] Mean sigma: 0.000000
[t-SNE] KL divergence after 250 iterations with early exaggeration: 70.454361
[t-SNE] KL divergence after 2000 iterations: 0.576950
(600, 2)

In [40]: output_notebook()
plot_tfidf = bp.figure(plot_width=700, plot_height=600, title="tf-idf clustering of t
tools="pan,wheel_zoom,box_zoom,reset,hover,previewsave",
x_axis_type=None, y_axis_type=None, min_border=1)

palette = d3['Category10'][len(tsne_tfidf_dfAff['Category'].unique())]
color_map = bmo.CategoricalColorMapper(factors=tsne_tfidf_dfAff['Category'].map(str).

plot_tfidf.scatter(x='x', y='y', color={'field': 'Category', 'transform': color_map},
legend='Category', source=tsne_tfidf_dfAff)
hover = plot_tfidf.select(dict(type=HoverTool))
hover.tooltips={"Reporter Name": "@ReporterName", "Category": "@Category", "Affiliations

show(plot_tfidf)

```

Observation: The bottom most cluster are associated to unknown attributes. And the right-most cluster belongs to PHDs and Stanford university.

## Finding Distorsion and Silhouette Scores for Affiliations

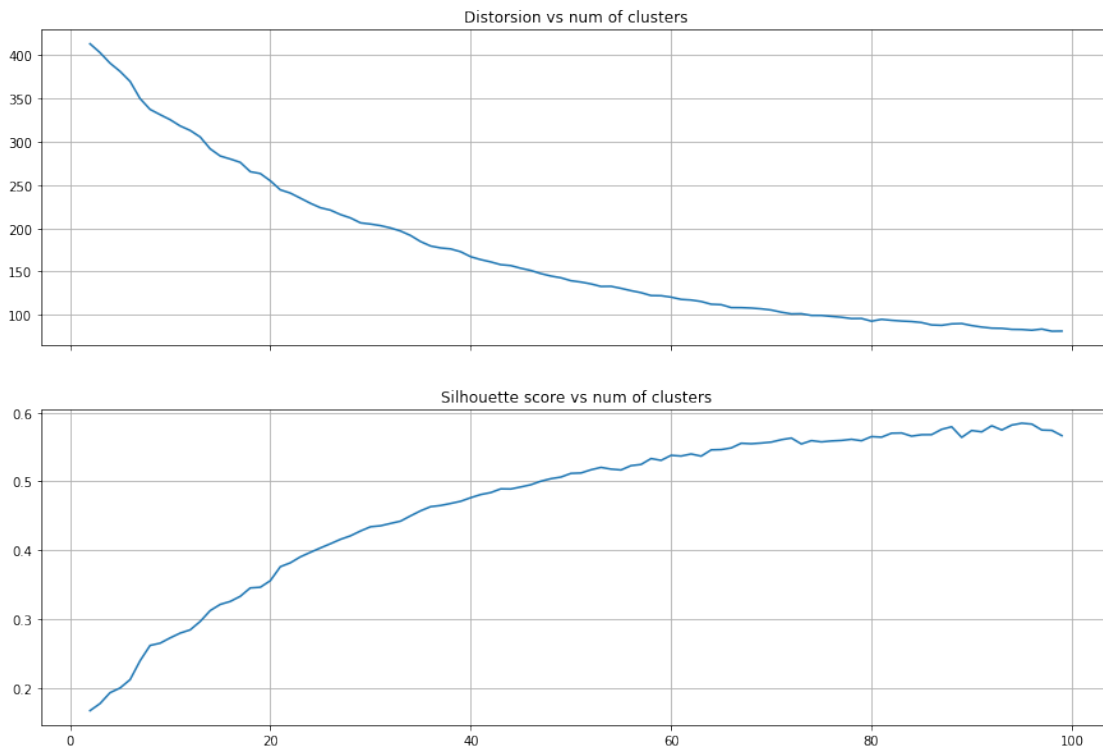
```
In [41]: distortions = []
        sil_scores = []
        k_max = 100
        for k in tqdm_notebook(range(2, k_max)):
            kmeans_model = MiniBatchKMeans(n_clusters=k, init='k-means++', n_init=1, random_s
                                     init_size=1000, verbose=False, max_iter=1000)
            kmeans_model.fit(vz1)
            sil_score = silhouette_score(vz1, kmeans_model.labels_)
            sil_scores.append(sil_score)
            distortions.append(kmeans_model.inertia_)

HBox(children=(IntProgress(value=0, max=98), HTML(value='')))
```

```
In [42]: f, (ax1, ax2) = plt.subplots(2, 1, sharex=True, figsize=(15, 10))

        ax1.plot(range(2, k_max), distortions)
        ax1.set_title('Distorsion vs num of clusters')
        ax1.grid(True)

        ax2.plot(range(2, k_max), sil_scores)
        ax2.set_title('Silhouette score vs num of clusters')
        ax2.grid(True)
```



There isn't an elbow curve which would tell us the optimal K. And since the error keeps on reducing it's difficult to know what's the correct K. Hence we have to choose the ratio of errors in the interval of K values. The error rate curve decreases less lower in between 20-40 as compared to the interval below 20. and hence k is selected as 20 to make sure there are good number of clusters as well.

```
In [43]: num_clusters = 20
         kmeans_model = MiniBatchKMeans(n_clusters=num_clusters, init='k-means++', n_init=1, r
                                     init_size=1000, batch_size=1000, verbose=False, max_iter=1000)
         kmeans = kmeans_model.fit(vz1)
         kmeans_clusters = kmeans.predict(vz1)
         kmeans_distances = kmeans.transform(vz1)
```

### 2.6.1 Affiliations : Clusters key words

```
In [44]: sorted_centroids = kmeans.cluster_centers_.argsort()[:, :-1]
         terms = vectorizerAff.get_feature_names()
         all_keywords = []
         for i in range(kmeans.n_clusters):
             topic_keywords = []
             for j in sorted_centroids[i, :10]:
                 topic_keywords.append(terms[j])
             all_keywords.append(topic_keywords)

         keywords_df = pd.DataFrame(index=['cluster_{0}'.format(i) for i in range(num_clusters)],
                                   columns=['keyword_{0}'.format(i) for i in range(10)],
                                   data=all_keywords)

         keywords_df
```

```
Out [44]:
```

	keyword_0	keyword_1	keyword_2 \
cluster_0	university	professor	state university
cluster_1	university	texas	york
cluster_2	universidad	del	researcher
cluster_3	affiliation	unknown affiliation	unknown
cluster_4	professor university	professor	university
cluster_5	center	medical center	medical
cluster_6	professor	university	york
cluster_7	scientist	senior	assistant
cluster_8	research fellow	fellow	research
cluster_9	hospital	university	oxford
cluster_10	science	technology	university
cluster_11	department	university	medicine
cluster_12	north carolina	carolina	north
cluster_13	institute	research	university
cluster_14	research	research associate	associate
cluster_15	studies	professor	university

cluster_16	newcastle	university	york
cluster_17	college	professor	hunter college
cluster_18	postdoctoral	fellow	university
cluster_19	associate	associate professor	professor

	keyword_3	keyword_4	keyword_5 \
cluster_0	state	school	engineering
cluster_1	harvard	doctoral	economics
cluster_2	professor	york	george
cluster_3	harvard	doctoral	economics
cluster_4	associate professor	associate	columbia
cluster_5	university	professor	columbia
cluster_6	harvard	economics	engineering
cluster_7	international	computer	science
cluster_8	university	hospital	newcastle
cluster_9	professor	department	center
cluster_10	professor	institute	engineering
cluster_11	professor	york	engineering
cluster_12	university north	university	anthropology
cluster_13	technology	new	biology
cluster_14	national	university	center
cluster_15	management	school	lecturer
cluster_16	harvard	economics	engineering
cluster_17	hunter	school	medicine
cluster_18	researcher	university oxford	cancer
cluster_19	university	state university	state

	keyword_6	keyword_7	keyword_8 \
cluster_0	laboratory	student	stanford
cluster_1	engineering	environmental	faculty
cluster_2	doctoral	economics	engineering
cluster_3	engineering	environmental	faculty
cluster_4	germany	michigan	colorado
cluster_5	harvard	washington	george
cluster_6	environmental	faculty	federal
cluster_7	university	york	florida
cluster_8	national	stanford university	postdoctoral
cluster_9	york	germany	engineering
cluster_10	phd	computer science	computer
cluster_11	mathematics	biology	university california
cluster_12	state university	state	environmental
cluster_13	director	oxford	professor
cluster_14	germany	medical	laboratory
cluster_15	international	assistant professor	assistant
cluster_16	environmental	faculty	federal
cluster_17	science	university	director
cluster_18	oxford	associate	center
cluster_19	physics	sciences	management

	keyword_9
cluster_0	phd
cluster_1	federal
cluster_2	environmental
cluster_3	federal
cluster_4	distinguished
cluster_5	national
cluster_6	fellow
cluster_7	economics
cluster_8	stanford
cluster_9	environmental
cluster_10	professor computer
cluster_11	california
cluster_12	federal
cluster_13	medicine
cluster_14	cancer
cluster_15	germany
cluster_16	fellow
cluster_17	economics
cluster_18	research
cluster_19	washington

```
In [45]: tsne_tfidf_dfAff['clusters'] = kmeans_clusters
         kmeans.n_clusters
```

```
Out[45]: 20
```

```
In [46]: kmeans_clusters1 = list(map(str,kmeans_clusters))
         tsne_tfidf_dfAff.head()
```

```
Out[46]:
```

	x	y	Category	ReporterName \
0	-9.728665	-1.908636	Magazine	diego bernardini
1	-18.391047	-1.789205	Magazine	raquel beer
2	37.898724	-33.167332	Newspaper	conrado morenoã
3	38.114857	-33.231438	Newspaper	rosario medina
4	36.284264	-32.078583	Newspaper	javier sampetro

	Affiliations	clusters
0	Instituto Universitario de Ciencias Medicas Fu...	0
1	Teacher Department Nursing Faculty Health Scie...	11
2	Profesor del Centro de Estudio de Tecnologias ...	2
3	Universidad del Rosario Facultad de Jurisprude...	2
4	Catedrático de ciencias del deporte	0

```
In [47]: color_lookup = dict(zip(set(kmeans_clusters),palette))
         tsne_tfidf_dfAff['colors'] = tsne_tfidf_dfAff['clusters'].apply(lambda x:color_lookup
```

```
In [48]: output_notebook()
         plot_tfidf = bp.figure(plot_width=700, plot_height=600, title="tf-idf clustering of t
```

```

tools="pan,wheel_zoom,box_zoom,reset,hover,previewsave",
x_axis_type=None, y_axis_type=None, min_border=1)

palette = d3['Category20'][kmeans.n_clusters]
color_map = bmo.CategoricalColorMapper(factors=list(set(kmeans_clusters1)), palette=p

plot_tfidf.scatter(x='x', y='y', color='colors',
                   legend='clusters', source=tsne_tfidf_dfAff)
hover = plot_tfidf.select(dict(type=HoverTool))
hover.tooltips={"Reporter Name": "@ReporterName", "Category": "@Category", "Cluster#": "@Cluster#"}

show(plot_tfidf)

```

## 2.6.2 Clustering on Interests

Next interesting attribute to look at is interest , although this doesn't hold that much significance in terms of reporting, it could still give us a fair idea of how similar the interests of reporters are.

### Preprocessing Text of Interests

```

In [49]: df['Interests'] = df['Interests'].apply(lambda x:toString(x))
         df['Cleaned Interests'] = df['Interests'].apply(lambda x:secondClean(x))
         vectorizer = TfidfVectorizer(min_df=5, analyzer='word', ngram_range=(1, 2), stop_words='english')
         vz2 = vectorizer.fit_transform(list(df['Cleaned Interests']))

         vz2.shape

Out[49]: (600, 105)

In [50]: svd = TruncatedSVD(n_components=50, random_state=0)
         svd_tfidf = svdAff.fit_transform(vz2)
         svd_tfidf.shape

Out[50]: (600, 50)

In [51]: tsne_model = TSNE(n_components=2, perplexity=10, verbose=1, random_state=0, n_iter=2000)
         tsne_tfidf = tsne_modelAff.fit_transform(svd_tfidf)
         print(tsne_tfidf.shape)

[t-SNE] Computing 31 nearest neighbors...
[t-SNE] Indexed 600 samples in 0.002s...
[t-SNE] Computed neighbors for 600 samples in 0.051s...
[t-SNE] Computed conditional probabilities for sample 600 / 600
[t-SNE] Mean sigma: 0.000000
[t-SNE] KL divergence after 250 iterations with early exaggeration: 75.665688
[t-SNE] KL divergence after 2000 iterations: 0.713738
(600, 2)

```

```

In [52]: tsne_tfidf_df1 = pd.DataFrame(tsne_tfidf)
        tsne_tfidf_df1.columns = ['x', 'y']
        tsne_tfidf_df1['Category'] = df['Category']
        tsne_tfidf_df1['ReporterName'] = df['Reporter Name']
        tsne_tfidf_df1['Interests'] = df['Interests']

In [53]: output_notebook()
        plot_tfidf = bp.figure(plot_width=700, plot_height=600, title="tf-idf clustering of t
            tools="pan,wheel_zoom,box_zoom,reset,hover,previewsave",
            x_axis_type=None, y_axis_type=None, min_border=1)

        palette = d3['Category10'][len(tsne_tfidf_df1['Category'].unique())]
        color_map = bmo.CategoricalColorMapper(factors=tsne_tfidf_df1['Category'].map(str).un

        plot_tfidf.scatter(x='x', y='y', color={'field': 'Category', 'transform': color_map},
            legend='Category', source=tsne_tfidf_df1)
        hover = plot_tfidf.select(dict(type=HoverTool))
        hover.tooltips={"Reporter Name": "@ReporterName", "Category": "@Category", "Interests":

        show(plot_tfidf)

```

## Distorsion and Silhouette Scores for Interests

```

In [54]: distortions = []
        sil_scores = []
        k_max = 100
        for k in tqdm_notebook(range(2, k_max)):
            kmeans_model = MiniBatchKMeans(n_clusters=k, init='k-means++', n_init=1, random_s
                init_size=1000, verbose=False, max_iter=1000)
            kmeans_model.fit(vz2)
            sil_score = silhouette_score(vz2, kmeans_model.labels_)
            sil_scores.append(sil_score)
            distortions.append(kmeans_model.inertia_)

        HBox(children=(IntProgress(value=0, max=98), HTML(value='')))

```


```

In [55]: f, (ax1, ax2) = plt.subplots(2, 1, sharex=True, figsize=(15, 10))

        ax1.plot(range(2, k_max), distortions)
        ax1.set_title('Distorsion vs num of clusters for interests')
        ax1.grid(True)

        ax2.plot(range(2, k_max), sil_scores)
        ax2.set_title('Silhouette score vs num of clusters for interests')
        ax2.grid(True)

```



## Clustering

```
In [56]: num_clusters = 20
         kmeans_model = MiniBatchKMeans(n_clusters=num_clusters, init='k-means++', n_init=1, random_state=42,
                                         init_size=1000, batch_size=1000, verbose=False, max_iter=1000)
         kmeans = kmeans_model.fit(vz2)
         kmeans_clusters = kmeans.predict(vz2)
         kmeans_distances = kmeans.transform(vz2)

In [57]: tsne_tfidf_df1['clusters'] = kmeans_clusters
         tsne_tfidf_df1['colors'] = tsne_tfidf_df1['clusters'].apply(lambda x: color_lookup.get(x))

In [58]: output_notebook()
         plot_tfidf = bp.figure(plot_width=700, plot_height=600, title="tf-idf clustering of topics",
                                tools="pan,wheel_zoom,box_zoom,reset,hover,previewsave",
                                x_axis_type=None, y_axis_type=None, min_border=1)

         palette = d3['Category20'][kmeans.n_clusters]
         color_map = bmo.CategoricalColorMapper(factors=list(set(kmeans_clusters)), palette=palette)

         plot_tfidf.scatter(x='x', y='y', color='colors',
                           legend='clusters', source=tsne_tfidf_df1)
         hover = plot_tfidf.select(dict(type=HoverTool))
```



```

hover.tooltips={"Reporter Name": "@ReporterName", "Category": "@Category", "Cluster#": "@Cluster#"}

show(plot_tfidf)

```

Here, the majority of interests are included in cluster 0 and small clusters are scattered all over the place. Thus no strong similarity in interests noted.

## Clusters of Keywords

```

In [59]: sorted_centroids = kmeans.cluster_centers_.argsort()[:, :-1]
terms = vectorizer.get_feature_names()
all_keywords = []
for i in range(kmeans.n_clusters):
    topic_keywords = []
    for j in sorted_centroids[i, :10]:
        topic_keywords.append(terms[j])
    all_keywords.append(topic_keywords)

keywords_df = pd.DataFrame(index=['cluster_{0}'.format(i) for i in range(num_clusters)],
                            columns=['keyword_{0}'.format(i) for i in range(10)],
                            data=all_keywords)

keywords_df

```

```

Out [59]:

```

	keyword_0	keyword_1	keyword_2 \
cluster_0	science	learning	social
cluster_1	nan	theory	design
cluster_2	engineering	software engineering	systems
cluster_3	health	public health	public
cluster_4	ecology	evolutionary	biology
cluster_5	cognitive	psychology	recognition
cluster_6	genetics	bioinformatics	genomics
cluster_7	conservation	ecology	biology
cluster_8	cancer	biology	clinical
cluster_9	medicine	care	nuclear
cluster_10	biology	molecular	structural
cluster_11	neuroscience	imaging	biology
cluster_12	modeling	science	care
cluster_13	psychology	social	media
cluster_14	theory	studies	development
cluster_15	management	capital	learning
cluster_16	education	sociology	medical
cluster_17	law	economics	management
cluster_18	climate	change	ecology
cluster_19	chemistry	materials	processing

	keyword_3	keyword_4	keyword_5 \
cluster_0	machine	physics	analysis
cluster_1	digital	dynamics	ecology

cluster_2	software	human	management
cluster_3	global	medical	psychology
cluster_4	evolution	theory	digital
cluster_5	neural	law	theory
cluster_6	gene	research	biology
cluster_7	conservation biology	animal	marine
cluster_8	physics	research	digital
cluster_9	physics	health	global
cluster_10	plant	bioinformatics	physiology
cluster_11	cognitive	theory	evolution
cluster_12	quality	computer	theory
cluster_13	clinical	climate	risk
cluster_14	physics	innovation	social
cluster_15	systems	intelligence	economics
cluster_16	biology	psychology	technology
cluster_17	regulation	health	theory
cluster_18	marine	conservation	gene
cluster_19	structural	physics	environmental

	keyword_6	keyword_7	keyword_8 \
cluster_0	machine learning	systems	computer
cluster_1	economics	education	energy
cluster_2	programming	information	interaction
cluster_3	care	human	research
cluster_4	dynamics	economics	education
cluster_5	environmental	digital	dynamics
cluster_6	computational biology	nuclear	computational
cluster_7	science	management	materials
cluster_8	dynamics	ecology	economics
cluster_9	science	environmental	development
cluster_10	marine	cancer	computational biology
cluster_11	dynamics	ecology	economics
cluster_12	epidemiology	digital	dynamics
cluster_13	behavior	management	economics
cluster_14	sociology	organizational	science
cluster_15	mining	public	ecology
cluster_16	research	marine	medicine
cluster_17	digital	dynamics	ecology
cluster_18	dynamics	economics	education
cluster_19	nuclear	biology	dynamics

	keyword_9
cluster_0	information
cluster_1	engineering
cluster_2	computer
cluster_3	cancer
cluster_4	energy
cluster_5	ecology

cluster_6	human
cluster_7	natural
cluster_8	education
cluster_9	digital
cluster_10	computational
cluster_11	education
cluster_12	ecology
cluster_13	education
cluster_14	neural
cluster_15	dynamics
cluster_16	science
cluster_17	education
cluster_18	energy
cluster_19	imaging

Cluster 0 is the biggest cluster but after looking the most frequent keyword it is still difficult to form a judgement on the interests.

## 2.7 Bigram analysis on Affiliations

Since with single words it is difficult to form an opinion , bigrams would help us to form a better inference of the common context in various texts.

```
In [60]: text = list(df['Cleaned Affiliations'])

In [61]: bigrams = [b for l in text for b in zip(l.split(" ")[:-1], l.split(" ")[1:])]

In [62]: from collections import Counter
counter = Counter()
for l in text:
    for t in range(0, len(l.split(" ")) - 1):
        counter.update([l.split(" ")[t] + " " + l.split(" ")[t+1]])
counter.most_common(10)

Out[62]: [('Unknown affiliation', 26),
          ('State University', 21),
          ('Associate Professor', 20),
          ('Assistant Professor', 11),
          ('Medical Center', 10),
          ('Professor University', 10),
          ('University North', 9),
          ('North Carolina', 8),
          ('University California', 8),
          ('Computer Science', 8)]
```

There are around 600 records and within these records only 26 Unknown affiliations occur> like wise with other bigrams, it could mean the affiliations are varied. Though, State University and Associate/Assistant Professors seem to lead in the data.

## 2.8 Bigram Analysis on Interests

```
In [63]: InterestText = df['Cleaned Interests'].str.lower()
InterestCounter = Counter()
for l in InterestText:
    for t in range(0, len(l.split(" "))-1):
        InterestCounter.update([l.split(" ")[t]+" "+l.split(" ")[t+1]])
InterestCounter.most_common(10)
```

```
Out[63]: [('machine learning', 14),
          ('artificial intelligence', 7),
          ('software engineering', 6),
          ('natural language', 6),
          ('public health', 5),
          ('data science', 5),
          ('language processing', 5),
          ('conservation biology', 5),
          ('computational biology', 5),
          ('molecular biology', 4)]
```

The counts though insignificant in comparison to the total data, but we can identify a general trend Interests have a bend towards scientific topics like Machine Learning ,Natural Language Processing and Biology.

## 2.9 Bigram Analysis on Publication Text

```
In [64]: PublicationCounter = Counter()
PublicationTexts = df['Cleaned Publications']
for l in tqdm_notebook(PublicationTexts):
    for t in range(0, len(l.split(" "))-1):
        PublicationCounter.update([l.split(" ")[t]+" "+l.split(" ")[t+1]])

HBox(children=(IntProgress(value=0, max=600), HTML(value='')))
```

```
In [65]: PublicationCounter.most_common(20)
```

```
Out[65]: [('united states', 136),
          ('present study', 82),
          ('results show', 81),
          ('paper presents', 77),
          ('climate change', 73),
          ('gene expression', 65),
          ('mental health', 64),
          ('risk factors', 61),
          ('case study', 59),
```

```
('health care', 57),  
('age years', 57),  
('purpose study', 55),  
('wide range', 55),  
('recent years', 53),  
('older adults', 53),  
('results indicate', 52),  
('data collected', 49),  
('aim study', 48),  
('little known', 48),  
('magnetic resonance', 47)]
```

No strong inference can be judged from here. Though, health and medicine seem to be a more commonly discussed subject.

In [ ]:

In [ ]: